KHUSHI SHARMA

Jaipur, India | +91 6378288504 | khushiveshnav123@gmail.com | LinkedIn | GitHub | Portfolio

SUMMARY

- Final-year B.Tech CSE student (CGPA 9.7/10) seeking an entry-level Applied ML / NLP role.
- Trained Transformer models for machine translation (BLEU 50 on 119K+ pairs) and business entity resolution (top 500 of 80,000+ in the Amazon ML Challenge 2026), deployed as Flask apps.
- Three data science internships covering forecasting, classification and automated reporting.

EDUCATION

Jaipur Engineering College (Affiliated to Rajasthan Technical University, Kota) | B.Tech CSE 2027
CGPA 9.7/10 (through Semester 6) | Ranked 2nd of 100 students in first year

SKILLS

ML & Statistics: Scikit-learn, XGBoost, Regression (Linear/Ridge/Lasso), Logistic Regression, Random Forest, Decision Trees, K-Means, SMOTE, Time-Series Forecasting, Anomaly Detection, Statistical Modeling, Model Validation, Benchmarking, Deep Learning, Neural Networks, TensorFlow, Keras, Transformers, GenAI / LLMs, RAG, Prompt Engineering, Reinforcement Learning

Data & Analytics: Python, SQL, PostgreSQL, Pandas, NumPy, Apache Spark, EDA, Feature Engineering, Statistical Sampling, Hypothesis Testing, Sensitivity Analysis, Data Modeling, Segmentation, KPI Analysis

Deployment: Flask, REST APIs, joblib, AWS (EC2/S3/Lambda), CI/CD, Git/GitHub, Vercel, React, HTML/CSS, JS (basic)

BI: Power BI, Tableau, DAX, Power Query, Excel, Matplotlib, Seaborn, Streamlit, Report Automation (python-docx)

CS Fundamentals: Data Structures & Algorithms, OOP, SDLC, Debugging, Smoke Testing, Technical Documentation

EXPERIENCE

Data Analytics with AI Intern | IBM (AICTE-approved) | GitHub Jun 2026 - Sep 2026

- Benchmarked Linear, Ridge and Lasso regression on a 1,000-record, 7-feature housing dataset with train-only scaling to avoid leakage; all three reached R^2 ~0.90 (CV R^2 0.894, MAE ~$16.1K), so regularization added no gain.
- Built an end-to-end house price prediction pipeline: Flask REST API (/predict, /schema, /health) and Streamlit UI deployed on Vercel, with automated retraining and Word/HTML reporting.

Data Science Intern | Xylofy AI | GitHub Apr 2026 - Jun 2026

- Built a Python pipeline and XGBoost 3-month sales forecast (time-based split) on 9,800 retail orders (2015-2018), with a team of three.
- Flagged regional sales swings (South -32% in 2016, Central +42% in 2017) and anomalous months with Isolation Forest; ~4-day shipping times were consistent across regions, pointing to market shifts rather than fulfillment delays.
- Grouped product sub-categories into 4 demand segments with K-Means and set a stocking strategy per segment; delivered an interactive Flask dashboard and an executive report with inventory recommendations.

Team Lead, Data Science & Analytics Intern | LinuxWorld Informatics (onsite) Jun 2025 - Aug 2025

- Raised classification F1-score by 12% on 50K+ records using statistical sampling and SMOTE.
- Deployed on AWS EC2 with CI/CD (S3, Lambda); automated 8 PostgreSQL pipelines (est. 156+ reporting hours saved yearly).
- Led a team of 3 that took AIW (All Is Well), a mood-aware conversation app, from idea to working product and final pitch at JAZBAA, LinuxWorld's student startup showcase, using Python, Flask, ML and DL.

PROJECTS

Amazon ML Challenge 2026: Business Entity Resolution (solo) | Python, Transformers, XGBoost

- Linked noisy, partial business records from three independent sources to a reference source in a 72-hour challenge, combining a Transformer with XGBoost; placed in the top 500 of 80,000+ registrants.

Neural Machine Translation Tool | Transformers, Flask | GitHub

- Trained a Transformer on 119K+ English-Spanish and English-French pairs (BLEU 50); deployed as a Flask web app with a GUI.

Mood Beats: Mood-Based Music Recommender | Python, Flask, React, Vite, Vercel | GitHub

- Trained a 4-class mood classifier (85% accuracy, saved with joblib); served via a Flask REST API and a React/Vite front end on Vercel.

Employee Attrition Prediction | Python, Pandas, Scikit-learn | GitHub

- Analyzed IBM HR data (1,470 employees, 16% attrition): leavers had a median income of $3,202 vs $5,204 for stayers; poor work-life balance roughly doubled attrition (31% vs 14%).
- Built a pipeline (80/20 split, SMOTE on training data) with Logistic Regression, Random Forest and Gradient Boosting; monthly income and overtime ranked as top drivers.

Healthcare Analytics Dashboard | Power BI, Tableau, DAX, Power Query | GitHub

- Designed a 3-page Power BI dashboard on a 1,000-patient dataset covering demographics, visits, insurance mix, cost KPIs (78.65% insurance coverage) and readmission ratio by hospital.

CERTIFICATIONS

- AWS Certified Cloud Practitioner
- Advanced SQL, HackerRank (2026)
- Google Analytics, Google (2025)
- Generative AI Essentials: Using LLMs to Work with Data, IBM (Aug 2026)
- Data Analytics, Deloitte (2026; hands-on Tableau work)
- Data Science & DSA, Learn and Build (2025)

ACHIEVEMENTS

- Amazon ML Challenge 2026: top 500 of 80,000+ as a solo participant; earned $250 in AWS credits.
- Core Member, Placement Cell: coordinated recruitment events and managed recruiter relations.
- Solved 70+ Data Structures & Algorithms problems on LeetCode.