Khushi Sharma

Data Scientist | Python, Machine Learning, Statistics, SQL
Jaipur, India | +91 6378288504 | Email | LinkedIn | GitHub | Portfolio

PROFESSIONAL SUMMARY

Entry-level Data Scientist and final-year B.Tech Computer Science student (CGPA 9.7/10) with three data science internships. Built an XGBoost time-series sales forecast with Isolation Forest anomaly detection and K-Means demand segmentation, and improved a classifier's F1-score by 12% on 50,000+ records using SMOTE. Experienced in feature engineering, model validation, hypothesis testing and communicating results to stakeholders. Top 500 of 80,000+ participants in the Amazon ML Challenge 2026.

TECHNICAL SKILLS

Machine Learning: Scikit-learn, XGBoost, Linear, Ridge and Lasso Regression, Logistic Regression, Random Forest, Gradient Boosting, Decision Trees, K-Means Clustering, SMOTE, Anomaly Detection, Isolation Forest, Model Validation, Cross-Validation, Benchmarking
Statistics: Hypothesis Testing, Statistical Sampling, Statistical Modeling, Time-Series Forecasting, Feature Engineering, Exploratory Data Analysis (EDA)
Deep Learning and GenAI: TensorFlow, Keras, PyTorch, Hugging Face Transformers, Neural Networks, NLP, GenAI, Large Language Models (LLMs)
Programming and Data: Python, SQL, PostgreSQL, Pandas, NumPy, Matplotlib, Seaborn
Deployment and Tools: Docker, Flask, REST APIs, Streamlit, AWS EC2, CI/CD, Git, GitHub, Vercel, joblib
BI: Power BI, Tableau, Excel

PROFESSIONAL EXPERIENCE

Data Analytics with AI Intern, IBM (AICTE-approved) | Jun 2026 – Sep 2026
- Benchmarked Linear, Ridge and Lasso regression on a 1,000-record, 7-feature housing dataset with train-only scaling to prevent data leakage; all three models matched (cross-validated R^2 0.894, MAE about $16.1K), so regularization added no gain.
- Built an end-to-end price prediction pipeline with a Flask REST API (/predict, /schema, /health), Streamlit interface on Vercel, automated retraining and Word/HTML reporting.

Data Science Intern, Xylofy AI | Apr 2026 – Jun 2026
- Built a Python pipeline and XGBoost 3-month sales forecast with a time-based train/test split on 9,800 retail orders (2015–2018).
- Detected regional sales anomalies (South -32% in 2016, Central +42% in 2017) with Isolation Forest; consistent shipping times of about 4 days showed market shifts rather than fulfillment delays.
- Segmented product sub-categories into 4 demand groups with K-Means and defined a stocking strategy per segment; delivered an interactive Flask dashboard and executive report.

Team Lead, Data Science and Analytics Intern, LinuxWorld Informatics | Jun 2025 – Aug 2025
- Improved classification F1-score by 12% on 50,000+ records using statistical sampling and SMOTE class-imbalance correction.
- Automated 8 PostgreSQL reporting pipelines on AWS EC2 with CI/CD, saving 156+ reporting hours per year.
- Led a team of 3 to take a mood-aware conversation app (Python, Flask, ML, deep learning) from idea to final pitch at the JAZBAA startup showcase.

PROJECTS

Amazon ML Challenge 2026: Business Entity Resolution *(Python, Transformers, XGBoost)*
- Linked noisy, partial business records from three independent sources to a reference source in 72 hours by combining a Transformer with XGBoost; ranked in the top 500 of 80,000+ registrants (solo).

Employee Attrition Prediction *(Python, Scikit-learn, SMOTE)*
- Analyzed IBM HR data (1,470 employees, 16% attrition); built Logistic Regression, Random Forest and Gradient Boosting models with an 80/20 split and SMOTE applied to training data only.
- Found monthly income and overtime to be top attrition drivers; poor work-life balance roughly doubled attrition (31% vs 14%).

MoodBeats: Mood-Based Music Recommender *(Python, Flask, React)*
- Trained a 4-class mood classifier (85% accuracy) and served it through a Flask REST API with a React/Vite front end on Vercel.

Neural Machine Translation Tool *(Transformers, TensorFlow, Flask)*
- Trained a Transformer on 119K+ English-Spanish and English-French sentence pairs (BLEU 50) and deployed it as a Flask web app.

Hugging Face Transformers Pipeline Toolkit *(Python, Hugging Face, PyTorch, Docker)*
- Benchmarked pretrained NLP models across 6 tasks with the Hugging Face pipeline API, including sentiment analysis (DistilBERT), zero-shot classification (BART-large-MNLI) and named entity recognition (BERT-large-cased).
- Configured GPT-2, DistilGPT2 and DistilRoBERTa to study text generation constraints and masked language modeling; containerized the toolkit with Docker for reproducibility.

EDUCATION

B.Tech, Computer Science and Engineering, Jaipur Engineering College (Rajasthan Technical University, Kota) | Expected May 2027
CGPA 9.7/10 | Ranked 2nd in first year

CERTIFICATIONS

AWS Certified Cloud Practitioner • Generative AI Essentials: Using LLMs to Work with Data, IBM (2026) • Advanced SQL, HackerRank (2026) • Data Analytics, Deloitte (2026) • Data Science and DSA, Learn and Build (2025)

ACHIEVEMENTS

- Top 500 of 80,000+ participants, Amazon ML Challenge 2026; earned $250 in AWS credits.
- Solved 70+ Data Structures and Algorithms problems on LeetCode.
- Core Member, Placement Cell: coordinated recruitment events and managed recruiter relations.